

MENG INDIVIDUAL PROJECT

DEPARTMENT OF COMPUTING

IMPERIAL COLLEGE OF SCIENCE, TECHNOLOGY AND MEDICINE

**Solbolt – a compiler explorer and
smart contract gas usage tracker for
Solidity**

Author:

Richard Jun Wei Xiong

Supervisor:

Dr. William Knottenbelt

May 26, 2022

Submitted in partial fulfillment of the requirements for the MEng Computing
(Artificial Intelligence and Machine Learning) of Imperial College London

Contents

1	Introduction	1
1.1	Overview	1
1.2	Planned contributions	3
1.3	Outline of report	3
2	Background	5
2.1	The Ethereum Blockchain	5
2.2	Smart contracts	5
2.3	The Ethereum Virtual Machine	6
2.3.1	Overview	6
2.3.2	Storage	6
2.3.3	Gas consumption	7
2.4	Solidity	8
2.4.1	The Solidity language	8
2.4.2	The Solidity compiler and EVM assembly	9
2.4.3	Yul	10
2.5	Mythril	11
2.5.1	Symbolic execution	11
2.5.2	Global states	12
2.5.3	Strategies	13
2.5.4	The LASER EVM	13
2.5.5	Plugins and Hooks	14
3	Related works	15
3.1	GASTAP	15
3.1.1	Summary	15
3.1.2	Results	16
3.1.3	Limitations	16
3.2	VisualGas	17
3.2.1	Summary	17
3.2.2	Results	17
3.2.3	Limitations	18
3.3	Summary of Research	18
4	Implementation	19

4.1	Feature summary	19
4.1.1	Solidity Compilation	19
4.1.2	Symbolic Execution	20
4.1.3	Gas Inspector	21
4.2	Architecture Design	21
4.2.1	UI Elements and Redux State Management	21
4.2.2	Monaco Editor	23
4.2.3	Etherscan Integration	24
4.2.4	Backend and REST API	24
4.2.5	Celery Task Queue	24
4.2.6	Deployment Pipeline	24
4.3	Mythril Extensions	25
4.3.1	Gas Meter Plugin and Gas Hooks	25
4.3.2	Function Tracker Plugin	27
4.3.3	Loop Gas Meter Plugin	29
4.3.4	Loop Mutation Detector Plugin	32
4.3.5	Removal of Z3 SMT Constraints	33
5	Evaluation	34
6	Ethical considerations	36
7	Project Plan	38

1

Introduction

Whereas most technologies tend to automate workers on the periphery doing menial tasks, blockchains automate away the center. Instead of putting the taxi driver out of a job, blockchain puts Uber out of a job and lets the taxi drivers work with the customer directly.

– Vitalik Buterin [1]

1.1 Overview

In recent years, Ethereum has emerged as a promising network for building next-generation, decentralised applications (DApps). Ethereum makes use of blockchain technology to achieve consensus in a distributed, public, and immutable ledger. It supports the running of arbitrary code, called smart contracts, in its own execution environment, called the Ethereum Virtual Machine (EVM). These smart contracts are typically written in a high level programming language, most notably Solidity, and are often used to power the transactional backends of various DApps. They are also used to define the behaviour of utility tokens or various non-fungible tokens (NFTs), which are simply smart contracts that extend a standard interface. Examples would be decentralised exchanges such as Uniswap [2] and Sushiswap, decentralised borrowing and lending platforms such as Aave [3] and Compound [4], algorithmic stablecoins such as Dai [5] and TerraUSD [6], and NFT-powered gaming ecosystems such as Axie Infinity [7] and the Sandbox [8].

Part of the appeal and increasing usage of smart contracts is that it allows transactions and exchanges to take place in a trustless, non-custodial, and cryptographically secure fashion. For example, in the case of a decentralised exchange, there is no need for two untrusted parties to hand over custody of their assets to a trusted intermediary, who would then conduct the actual exchange. Smart contracts are instead able to conduct the exchange of tokens or assets in a single, atomic transaction, which will either succeed, or revert and leave both parties unaffected. This powerful ability

is imparted by Ethereum’s immutable and decentralised nature, as no single party can control or alter the state of the blockchain directly.

The result is an explosion in smart contract development and total value locked (TVL), from around \$600 million to over \$150 billion in the span of two years [9]. However, this has also brought on scalability issues that limit the overall usefulness of the Ethereum network. In simple terms, smart contract transactions each require a certain amount of “gas” in order to be processed, depending on its complexity. Users bid for the price of gas they are willing to pay, in ether, and miners typically choose to process the transactions with higher bids. Since only a limited number of transactions can be processed in each block, this results in an efficient market economy between miners and users.

However, when more users start utilising the network, the gas prices naturally start rising as demand grows. Lately, due to the boom in decentralised finance (DeFi) platforms and NFTs, the gas price has surged by almost 10 times its price compared to two years ago [10]. Transactions are costing up to \$63 on average in November [11], and this has made the Ethereum blockchain inhibitive to use for most users [12]. Uniswap, a popular decentralised exchange for trading tokens on the Ethereum network, has spent more than \$8.6 million per day in gas fees alone, as reported last November [13]. Therefore, there is a clear need not only for better scaling of the Ethereum blockchain, but also for developers creating DApps on the Ethereum blockchain to optimise their smart contract code for gas usage.

Despite the enormous cost savings that can be achieved via more gas-efficient smart contracts, most of the development on Ethereum smart contract analysis currently surrounds verification of smart contracts rather than gas optimisation [14]. The Solidity compiler itself does offer an estimate for gas usage [15], although only at a course-grained function-call level, making analysis of code within each function difficult. It also only performs a rudimentary best-effort estimate, and outputs infinity whenever the function becomes too complex. Given that Solidity is a language under active development with numerous breaking changes implemented at each major version release [16] [17], the few tools designed with gas estimation in mind are also either no longer maintained or broken, or inadequate for fine-grained, line-level gas analysis. We will discuss some of these existing tools in the background section later.

Currently, most Solidity developers simply make use of general tools such as the Remix IDE, Truffle and Ganache for smart contract development [18]. Such tools are not specialised for detailed gas usage profiling, and developers may potentially miss out on optimisations that may save up to thousands of dollars worth of gas [19], or worse, possibly render a contract call unusable due to reaching the gas limits of each block [20]. Gas optimisation is a step that every developer has to carefully consider when deploying an immutable contract on the blockchain, and therefore creating a tool that calculates the gas usage for each bytecode instruction and line of code could greatly enhance and simplify such analysis, not only for smart contracts on the Ethereum blockchain, but also for all EVM compatible chains such as the

Binance Smart Chain, Polygon, and the Avalanche Contract Chain [21].

1.2 Planned contributions

With the goal of gas optimisation in mind, this project aims to improve on current smart contract development tools by creating a smart contract compiler explorer tool, which can also display a heatmap of the gas usage of each bytecode instruction, while mapping them to the corresponding lines in the user's source code. This is inspired by Godbolt [22], a popular compiler explorer tool for C++ for finding optimisations in the GCC compiler. The following summarises the intended contributions this project plans to achieve.

- 1. Compilation and mapping of each line of user Solidity code into their corresponding EVM bytecode and Yul intermediate representation.**

Similar to Godbolt which is able to easily map and visualise which assembly instructions correspond to its high-level C++ code, we want to achieve a similar result with EVM assembly bytecode and its corresponding Solidity code. We also want to be able to generate the Yul intermediate representation for the Solidity code, which is a recently introduced language between Solidity and EVM assembly.

- 2. Analysis of gas consumption via symbolic execution for each EVM instruction**

The schedule for gas consumption of EVM instructions is carefully laid out in the Ethereum Yellow Paper [23]. Using this, we want to conduct an analysis of gas consumption for each line of Solidity code, instead of simply at a function call level. To do this, we propose using symbolic execution to traverse each possible execution path, and then aggregating the gas costs for each basic block. To address the problem of path explosion, we also propose the implementation of a non-iterative loop bounds analysis that can generate a parametric loop bound using polytopes. Then, we can allow users to define these external parameters to obtain a more precise estimate.

- 3. Generation of heatmaps of gas usage per line of code**

To make visualising gas usage simpler, we plan on generating a heatmap on top of the direct mapping from EVM bytecode instructions to Solidity code, which would display at a glance which lines of code are consuming the most amount of gas, and may potentially have room for optimisation.

1.3 Outline of report

In Chapter 2, we provide a summary of smart contracts on Ethereum, the internals of the Ethereum Virtual Machine and its gas schedule, as well as a background on Solidity and the Yul intermediate representation. In Chapter ??, we describe our

proposed implementation to estimate gas costs using symbolic execution, and a non-iterative algorithm using polytopes to statically infer loop bounds. In Chapter 3, we briefly examine two related projects, as well as their results and limitations. In Chapter 5, we propose the method we intend to measure the project’s quality and effectiveness by. In Chapter 6, we examine the possible ethical issues that may arise from this project. Finally, in Chapter 7, we discuss the proposed timeline for the completion of this project, including stretch goals and fallback goals under various circumstances.

2

Background

2.1 The Ethereum Blockchain

The Ethereum Blockchain can be described as a decentralised, transaction-based state machine [23]. It runs on a proof-of-work system, where, simply put, miners run a computationally difficult algorithm, called Ethash [24], to find a valid nonce for a given block through trial and error. Once a valid nonce is generated, it is very easy for other clients to verify, but almost impossible to tamper with, since changing even one transaction will lead to a completely different hash [25]. The highly decentralised nature of the network makes it incredibly costly and difficult to issue malicious blocks or reorder previous blocks.

The world state of the Ethereum blockchain is a mapping between 160-bit addresses and the corresponding account state. This world state can be altered through the execution of transactions, which represent valid state transitions. Formally:

$$\sigma_{t+1} \equiv \Upsilon(\sigma_t, T)$$

where Υ is the Ethereum state transition function, T is a valid transaction, and σ_t and σ_{t+1} are neighbouring states [23]. The main innovation of Ethereum is that Υ allows any arbitrary computation to be executed, and σ is able to store any arbitrary state, not limited to just account balances. This leads us to the concept of smart contracts, which are arbitrary code deployed and stored on the Ethereum blockchain, that can be triggered by contract calls in the form of transactions submitted to the blockchain.

2.2 Smart contracts

The idea of smart contracts was first introduced in 1994 by Nick Szabo [26], where he describes it as a "computerised transaction protocol that executes the terms of a contract". Contractual agreements between parties can be directly embedded within the systems that we interact with, which are able to self-enforce the terms of the contract in a way that "minimise[s] exceptions both malicious and accidental, and

minimise[s] the need for trusted intermediaries”, akin to a digital vending machine [27].

The Ethereum blockchain is one of the first widely successful and adopted execution environment for smart contracts. Due to the decentralised and tamper-resistant properties of Ethereum, smart contracts deployed on Ethereum can operate as autonomous actors [28], whose behaviours are completely deterministic and verifiable. Christidis et al also likens them to stored procedures in a relational database [28].

It is important to note that smart contracts must be completely deterministic in nature, or else each node in the decentralised network will output different resulting (but valid) states.

2.3 The Ethereum Virtual Machine

2.3.1 Overview

Smart contracts on Ethereum run in a Turing-complete execution environment called the Ethereum Virtual Machine (EVM). The EVM runs in a sandboxed environment, and has no access to the underlying filesystem or network. It is a stack machine that has three types of storage – storage, memory, and stack [29], explained in detail in section 2.3.2.

The EVM runs compiled smart contracts in the form of EVM opcodes [30], which can perform the usual arithmetic, logical and stack operations such as XOR, ADD, and PUSH. It also contains blockchain-specific opcodes, such as BALANCE, which returns address balance in wei.

An EVM transaction is a message sent from one account to another, and may include any arbitrary binary data (called the *payload*) and Ether [30]. If the target account is a *contract account* (meaning it also stores code deployed on the blockchain), then that code is executed, and the payload is taken as input.

The EVM also supports message calls, which allows contracts to call other contracts or send Ether to non-contract accounts [30]. These are similar to transactions, as they each have their own source, target, payload, Ether, gas and return data.

2.3.2 Storage

Every account contains a *storageRoot*, which is a hash of the root node of a Merkle Patricia tree that stores the *storage* content of that account [23]. This consists of a key-value store that maps 256-bit words to 256-bit words [30]. This storage type is the most expensive to read, and even more costly to write to, but is the only storage type that is persisted between transactions. Therefore, developers typically try to reduce the amount of storage content used, and often simply store hashes to off-chain storage solutions such as InterPlanetary File System [31] or Arweave permanent storage [32]. A contract can only access its own storage content.

Memory is the second data area where a contract can store non-persistent data. It is byte-addressable and linear, but reads are limited to a width of 256 bits, while writes can be either 8-bit or 256-bits wide. [30] Memory is expanded each time an untouched 256-bit memory word is accessed, and the corresponding gas cost is paid upfront. This gas cost increases quadratically as the memory space accessed grows.

The *stack* is the last data area for storing data currently being operated on, since the EVM is stack-based and not register-based [30]. This has a maximum size of 1024 elements of 256-bit words, and is the cheapest of the three types to access in general. It is possible to copy one of the topmost 16 elements to the top of the stack, or swap the topmost element with one of the 16 elements below it. Other operations (such as ADD or SUB) take the top two elements as input, and push the result on the top of the stack. It is otherwise not possible to access elements deeper within the stack without first removing the top elements, or without moving elements into memory or storage first.

2.3.3 Gas consumption

Gas refers to the "the unit that measures the amount of computational effort required to execute specific operations on the Ethereum network" [33]. It is the fee that is paid in Ether in order to successfully submit and execute a transaction on the blockchain. This mechanism is introduced in order to avoid malicious actors from spamming the network, either by submitting many transactions at once in a denial of service attack, or by running accidental or intentional infinite loops in smart contract code. It also incentivises miners to participate in the network, as part of the gas fees (in the form of tips as introduced in EIP-1559 [34]) is given to the miner .

Each block also has a block limit, which specifies the maximum amount of computation that can be done within each block. This is currently set at 30 million gas units, [33]. Transactions requiring more gas than the block limit will therefore always revert, which means that all executions will either eventually halt, or hit the block limit and revert.

The schedule for how gas units are calculated is described in detail in Appendix G and H of the Ethereum Yellow Paper [23]. In summary, each transaction first will require a base fee of 21000 gas units, which "covers the cost of an elliptic curve operation to recover the sender address from the signature as well as the disk and bandwidth space of storing the transaction" [35]. Then, depending on the opcodes of the contract being executed, additional units of gas will be used up. Most opcodes require a fixed gas unit, or a fixed gas unit per byte of data. An exception to this is the gas used for memory accesses, where it is calculated as such [23]:

$$C_{mem}(a) \equiv G_{memory} \cdot a + \left\lfloor \frac{a^2}{512} \right\rfloor$$

Here, G_{memory} is the gas unit paid for every every additional word when expanding memory. a is the number of memory words such that all accesses reference valid

memory. Therefore, memory accesses are only linear up to 724B, after which it becomes quadratic and each memory expansion costs more.

Another exception would be the SSTORE and SELFDESTRUCT instructions. For these instructions, the schedule is defined as follows [23]:

Name	Value	Description
G_{sset}	20000	Paid for an SSTORE operation when the value of the storage bit is set to non-zero from zero.
G_{reset}	2900	Paid for an SSTORE operation when the value of the storage bit's zeroness remains unchanged or is set to zero.
R_{sclear}	15000	Refund given (added into a <i>refund counter</i>) when the storage value is set to zero from non-zero.
$G_{\text{selfdestruct}}$	5000	Amount of gas to pay for a SELFDESTRUCT operation.
$R_{\text{selfdestruct}}$	24000	Refund given (added into a <i>refund counter</i>) for self-destructing an account.

The refund counter tracks the units of gas that is refunded to the user upon successful completion of a transaction, which means that the transaction does not revert or invoke an out-of-gas exception. Only up to half of the total gas used by the transaction can be refunded using the refund counter. This mechanism is introduced to encourage freeing up storage resources used by the Ethereum blockchain.

Therefore, we can see that the estimation of gas usage for any given function is not trivial or easily predictable, and often depends on the current state of the blockchain, as well as the input parameters given. It would also require detailed modelling of the gas and refund counter, as well as other internals of the EVM implementation.

2.4 Solidity

2.4.1 The Solidity language

Solidity is a statically-typed, object-oriented, high-level programming language for developing Ethereum smart contracts [36]. It is a language in active development with numerous breaking changes in each major release, all of which are documented extensively online [15]. We will summarise some of its notable features in this section.

Inheritance. Solidity supports inheritance, by extending other contracts. These contracts can be interfaces (*abstract contracts*), with incomplete implementations of function signatures. A prime example of this would be the ERC-20 and ERC-721 interfaces [37] [38], which are standard interfaces that Ethereum tokens and non-fungible tokens respectively are expected to follow.

Libraries. Solidity also supports the use of libraries, which can contain re-usable functions or structs that are referenced by other contracts. OpenZeppelin for example has a wide range of utility libraries for implementing enumerable sets, safe math, and so on [39].

Function scopes. Solidity allows developers to define the scopes of each function – internal, external, private, or public [40]. internal functions can only be called by the current contract or any contracts that extend it, and are translated into simple jumps within the EVM. The current memory is not cleared, making such function calls very efficient. private functions form a subset of internal functions, in that it is stricter by only being visible to the current contract. external functions can only be called via transactions or message calls, and all arguments must be copied into memory first. This means that they can only be called "internally" by the same contract via the `this` keyword, which effectively makes an external message call to itself. public functions are a superset of external functions, in that they can be called internally or externally.

Storage types. Solidity supports the explicit definition of where variables are stored – storage, memory, calldata, or stack. calldata is an immutable, non-persistent area for storing *function arguments* only, and is recommended to be used whenever possible as it avoids copies or modification of data. storage and memory stores the variable in the respective data areas, as described in section 2.3.2.

2.4.2 The Solidity compiler and EVM assembly

The Solidity compiler translates the high-level Solidity code into low-level EVM bytecode, as well as generating other metadata such as the contract Application Binary Interface (ABI), and coarse gas estimates for each function. It supports a built-in optimiser, which takes an `--optimize-runs` parameter. It may be interesting for developers to observe what differences this parameter makes within the bytecode.

To illustrate the operation of the compiler, let us examine the following example contract, as well as the (truncated) output from the Solidity compiler.

```

1 pragma solidity >=0.5.0 <0.9.0;
2 contract C {
3     function one() public pure returns (uint) {
4         return 1;
5     }
6 }
```

Listing 2.1: An example Solidity contract

```

1 ===== test.sol:C =====
2 ...
3 sub_0: assembly {
4     /* "test.sol":33:123  contract C {... */}
5     mstore(0x40, 0x80)
6     callvalue
7 ...
8     tag_1:
9     pop
10    jumpi(tag_2, lt(calldatasize, 0x04))
```

```

11     shr(0xe0, calldataload(0x00))
12 ...
13     tag_5:
14         /* "test.sol":87:91  uint */
15     0x00
16         /* "test.sol":111:112  1 */
17     0x01
18         /* "test.sol":104:112  return 1 */
19     swap1
20     pop
21         /* "test.sol":51:120  function one() public pure returns (uint) {... */
22     swap1
23     jump // out
24         /* "#utility.yul":7:84    */
25 ...
26         /* "#utility.yul":7:84    */
27     tag_9:
28         /* "#utility.yul":44:51    */
29 ...
30
31     auxdata: [...]
32 }

```

Listing 2.2: EVM assembly from the Solidity compiler

As seen in Listing 2.2, the Solidity code is translated into subroutines with multiple tags as jump destinations, similar to the GCC compiler. It also includes control flow instructions such as JUMPI and JUMP. Since the EVM is a stack machine, each JUMPI and JUMP instruction does not have a corresponding target as an argument. Instead, the jump target is specified by the value at the top of the stack, which makes control flow analysis somewhat more complicated. Valid jump addresses are also specified via the JUMPDEST instruction. Debug information is embedded into the output, which describes the code fragment that a particular segment of assembly maps to, similar to the -g flag in the GCC compiler.

An additional `#utility.yul` file is also generated, which contains extra Yul code (see Section 2.4.3) that are created automatically by the compiler, and its addition is triggered by the usage of specific language features, such as the ABI encoder for function parameters and return values. They are not specific to the user's source code, but are likely to be executed in control flow jumps from other parts of user code.

2.4.3 Yul

Yul is an intermediate language that is designed to be compiled into bytecode for different backends [41], such as future EVM versions and EWASM, and acts as a common denominator for all current and future platforms. Yul is also designed to be human readable, with high level constructs such as `for` and `switch`, but still suitable for whole-program optimisation. Performing analysis on the Yul IR also has the advantage over Solidity code that it reveals certain "*hidden loops*" when dealing with variable sized data types like `string` and `bytes`.

Due to the additional information that Yul provides in terms of control flow and stack allocation, the project originally planned to perform static loop analysis on the generated Yul intermediate representation. This idea was, however, abandoned for the following reasons:

1. Yul is only released recently in Solidity compiler v0.7.5 and above, and this severely limited the dataset of contracts available for evaluation purposes.
2. Performing the static analysis required some additional instrumentation on the compiler level, and in order to support all the compiler versions required for evaluation, one has to modify and recompiler every compiler version between v0.7.5 and v0.8.14 (the latest as of today).
3. Yul is still an experimental feature, and will likely be subject to frequent changes in terms of its language semantics. As such, it is possible that any static analysis performed using the current state of Yul may be rendered obsolete in the near future. Performing analysis on EVM assembly however, is more likely to still remain relevant for the time being, as it is the fundamental machine code that the Ethereum Virtual Machine runs on, and any changes to it will require a hard fork of the entire blockchain.

As such, instead of performing loop bounds inference on the Yul intermediate representation, the project has instead performed loop detection on the EVM assembly, and calculated per iteration gas costs on any possible loops, which will be explained in detail later in the implementation section.

2.5 Mythril

2.5.1 Symbolic execution

Symbolic execution is a technique of programme analysis that involves the execution of a programme using symbolic inputs rather than concrete inputs [42]. Then, it proceeds down each possible execution path if the conditions for it are satisfiable using a SMT solver, in order to determine some security properties of a programme, such as whether a particular code path is reachable or not, or if a certain assertion can be violated during the execution.

Compared to other techniques such as dynamic analysis, which evaluates a programme input-by-input rather than path-by-path, symbolic execution may provide increased efficiency and coverage, especially if the input space is large. This may be the case for Ethereum smart contracts, where inputs are often 256-bit unsigned integers, user-defined structs, or even dynamic arrays of strings and bytes. Analysing a smart contract via dynamic analysis may require thousands of concrete transactions to achieve a desirable level of coverage [43], while symbolic analysis may only require traversing a few hundred paths based on our testing.

A common problem when using symbolic execution for programme analysis is path explosion. The number of possible execution paths increases exponentially with an increase in code size, and can possibly be infinite in the case of unbounded loops. However, due to having to meet the gas limit of Ethereum blocks, smart contracts are typically much simpler compared to executable desktop programmes. The Ethereum Virtual Machine is also much easier to model due to its deterministic and sequen-

tial nature with a limited number of opcodes, unlike a traditional operating system with its complex system calls (which may also be asynchronous in nature) and concurrency handling. As such, these observations make symbolic execution a suitable analysis method for smart contract gas analysis, and we later find that even using a simple bounded loop strategy allows us to achieve a very high coverage for most smart contracts evaluated.

There are several symbolic execution engines that already exist for the Ethereum Virtual Machine. For example, tools such as Oyente [44] (and EthIR, an extension of it,) are able to perform symbolic execution and produce a control flow graph of the basic blocks executed. However, when testing their capabilities, it was found that they were no longer actively maintained, and contained many errors that required patching before they worked. They also did not support the latest Solidity version, and were missing numerous newly introduced opcodes, such as SHR (shift right) and SHL (shift left). There are also other libraries such as Manticore [45], but they were poorly documented and difficult to extend.

Ultimately, the symbolic execution engine chosen for the project is Mythril [46], which is a symbolic execution engine originally designed by ConsenSys for the verification of smart contracts and analysis of possible path conditions that are able to induce undesired behaviour, such as killing a smart contract. This project then extended Mythril's original capabilities, making use of its detailed of the EVM and how it handles each opcode, to track gas usage along each execution path and for each piece of code.

As a large part of the work done by this project is on directly extending the Mythril execution engine, it is helpful to understand how its internals work.

2.5.2 Global states

Mythril keeps track of the different program execution states via "global states", each of which is a complete snapshot of the Ethereum virtual machine at a given program point. Each global state contains the following key information:

1. World State

The World State, as described in the Ethereum yellow paper, contains information about each account and their balances, as well as their storage state and bytecode if they are smart contract accounts. It also keeps track of the current transaction sequence.

2. Environment

The Environment keeps track of the current execution environment for the symbolic executor, and contains the current active smart contract account, symbolic or concrete values for the sender, origin and calldata, as well as values for the gas price.

3. Machine State

The Machine State, represented by μ in the Ethereum yellow paper, keeps track of the current internal state of the Ethereum Virtual Machine, such as the stack, the programme counter, the state of the memory, as well as the amount of gas used so far.

4. Annotations

Annotations can also be attached to the global state, which are objects that can be accessed and modified by custom plugins to keep track of information relevant to the plugin. Plugins are passed from one global state to the next after performing the current opcode, and reset with each new transaction. These are used extensively by this project for extending Mythril with custom gas analysis plugins.

2.5.3 Strategies

Given a work list of global states to explore and execute, the symbolic executor can employ the following strategies in order to decide which paths to execute first:

1. **Breadth First Search:** Selects the next state from the front of the work list
2. **Depth First Search:** Selects the next state from the end of the work list
3. **Naive Random Search:** Selects the next state randomly from the work list with equal probability
4. **Weighted Random Search:** Selects the next state randomly from the work list with probability inversely proportional to the depth of that state within the list

After selecting the next state and executing its current opcode, the symbolic executor may obtain zero new states (if a REVERT, STOP or RETURN was executed), one new global states (for most opcodes), or multiple global states (if there was a conditional branch). It then appends the new states (if any) to the back of the work list, and repeats the process until the work list is empty.

Each strategy can also be extended with the bounded loops strategy (which stops execution once a given loop bound is reached), or call depth limit strategy (which limits the depth of the call stack).

2.5.4 The LASER EVM

The entire symbolic execution pipeline is then implemented as the LASER EVM. A UML diagram describing how the pipeline works can be found in Appendix A. In summary, the LASER EVM proceeds via the following steps:

1. The symbolic executor first instruments the LASER EVM with the necessary plugins, and initialises the world state, either with symbolic values or with user-configured concrete values.

2. The LASER EVM sends a contract creation transaction, which creates the active smart contract account that we want to perform our analysis on. This also returns a list of initial open global states, which the symbolic executor will explore.
3. The LASER EVM next begins executing the first transaction on the contract that was created. It selects the first global state to execute using the current strategy, and then calls the function responsible for modelling the current opcode. One or more new global states are then created from the current global state, and appended to the end of the work list.
4. If STOP or RETURN was reached, a transaction end signal is sent, and a new open global state is created by copying the previous world state of the ending transaction. These will be executed in the next transaction. If REVERT was reached, all changes are instead reverted.
5. Execution continues until the work list is empty. Then, the entire process repeats with the new open global states being executed in the next transaction, until the transaction limit defined by the user is reached.

2.5.5 Plugins and Hooks

The LASER EVM can be easily extended with additional functionality using Mythril's plugin and hooks system. Additional custom plugins can be instrumented into the LASER EVM, by registering certain functions to be called by hooks. These hooks are then invoked at several particular points of execution, such as:

- At the start and end of the entire symbolic execution pipeline
- At the start and end of each symbolic transaction
- When a new global state is being added or executed
- During the execution of an opcode (but before the new global state is returned), for a particular opcode or for all opcodes
- Before or after the execution of an opcode (with the initial global state or new global states returned), for a particular opcode or for all opcodes

Each hook, when invoked, calls all the functions that are registered with it, and usually also passes the current global state or list of global states, depending on its type. The functions (which are part of a plugin) can then make use of custom annotations within the global state to keep track of relevant properties about the current execution state or perform custom tasks.

This project makes extensive use of the modularity and ease of customisation provided by this plugin system, in order to extend Mythril's capabilities for gas cost analysis, as will be explained in chapter 4.

3

Related works

Most Solidity analysis tools currently still revolve around programme verification, such as Oyente, Mythril, and Manticore. These tools focus on analysing reachable vulnerable states within a smart contract, and checks for common exploits such as reentrancy and unprotected self destructs, and ultimately differ from the objectives of this project. MadMax is a verification tool that focuses on detecting states that may result in an out of gas error, but also does not focus on calculating gas bounds for smart contract functions, and therefore remains out of scope. Here, we will discuss two related projects, GASTAP and VisualGas, which both focus on a similar objective of estimating gas costs with two different approaches.

3.1 GASTAP

3.1.1 Summary

GASTAP is a tool developed by Albert et al [47] that claims to be able to infer parametric gas bounds for smart contracts. These inferred bounds may depend on the function arguments provided, or on the blockchain and smart contract state. GASTAP conducts its analysis on the EVM bytecode level, for which the schedule of gas fees is defined.

GASTAP achieves this in 3 stages. First, it generates a *stack-sensitive control flow graph (S-CFG)* for the EVM bytecode. To do this, it first calculates the set of basic blocks in an EVM programme, and detects the set of valid jump destinations (marked by the JUMPDEST instruction). Then, they proceed via symbolic execution for each block to produce the stack state after executing each EVM instruction. The key idea here is that since JUMP instructions jump to the programme counter at the top of the stack, control flow when reaching a basic block therefore then depends on the set of stack items at that point which hold valid jump destinations. Thus, if two execution traces reach the same basic block with the same set and locations of jump destinations in their stacks, their nodes in the control flow graph can be safely merged. Else, a new node for the new stack state is pushed into the control flow

graph, and the process repeats iteratively until the graph is consistent. Their earlier paper also provides a proof for how this model is sound and over-approximates the jumping information for a programme [48]. The tool they used to perform this step – EthIR, an extension of the Oyente tool – is also open source and available on their GitHub repository.

Next, GASTAP produces a *rule-based representation (GAS-RBR)* from the S-CFG generated in the previous step. Each for each block, a new rule is constructed, with the edges indicating invocations of the generated rules. The original EVM instructions are wrapped in a *nop* functor within the GAS-RBR, to allow for precise calculation of the gas costs later on. This representation essentially abstracts the relationships between different items within the stack, as well as the boolean guards involved involved in making conditional jumps. Then, in order to calculate gas bounds, GASTAP also defines a set of semantics for the GAS-RBR, and separates the gas costs into two parts – the *opcode cost* and the *memory cost*. The opcode semantics takes the EVM opcode, the state of the stack, and the mapping of state variables when reading the opcode to return the corresponding gas cost of the instruction. The memory semantics similarly returns the highest memory slot used by an EVM instruction, since memory gas cost is only determined by each expansion of memory slot as introduced in Section 2.3.3.

Finally, the current GAS-RBR allows for the calculation of gas costs for concrete executions, for which all parameters involved are known concretely. However, in order to calculate parametric bounds, GASTAP first transforms the local memory and storage accesses by the GAS-RBR into local variables. Then, it makes use of a resource analysis tool called SACO to solve the constraints set by the GAS-RBR, and output a final parametric gas cost of a function call. This is the main contribution of this project, as there is no other tool available to our knowledge that is able to produce a parametric bound on gas costs.

3.1.2 Results

GASTAP is then evaluated against 34,000 Ethereum smart contracts, which was completed in 407.5 hours, or around 16 days of execution time. It was also found that GASTAP had a failure rate of 0.85%. Then, the gas bounds generated by GASTAP was evaluated against 4000 transactions of 300 top-valued Ether smart contracts. The average precision obtained was not published, but the overhead calculated from GASTAP was about 10% to 50% of the real transaction gas costs, with some overheads of up to 600% recorded. The authors claim this is because GASTAP has to take into account all possible input values, in order to calculate the worst case bounds.

3.1.3 Limitations

Some limitations of the GASTAP tool is that it only calculates gas bounds on a function-call level, similar to the Solidity compiler but with possibly greater precision. Our project aims to improve on this by calculating gas used for each line of

Solidity code.

In addition, the GASTAP demo is no longer working or being maintained, and their latest supported version of Solidity is 0.7.1, which was released in September 2020. In addition, the SACO analyser used was developed by themselves and is closed-source, with no available binaries to the public, and we could not verify their claims.

3.2 VisualGas

3.2.1 Summary

VisualGas is a tool developed by Signer [43] with a similar goal for allowing developers to visualise the gas costs incurred by each line of code, and test best and worst case executions before deployment. However, in contrast to the static resource analysis of GASTAP, and our proposed implementation of symbolic execution, VisualGas makes use of dynamic programme analysis to estimate gas bounds of each instruction. [43] claims that this method is more precise than static analysis, although may possibly lead to less coverage.

VisualGas was implemented in 3 parts. First, since dynamic analysis is performed, VisualGas needs a way to collect traces of executions of a given smart contract. It uses a Go-Ethereum client locally to execute transactions, and collects traces in the form of what instruction was executed at which programme counter (PC), the remaining amount of gas, the gas cost of each step, the call stack depth and memory state, as well as any storage slots accessed at each step. This trace collection is repeated for every new test input state that is generated. Then, it uses the source mapping of build artifacts generated by Truffle, and processes it to map programme counters to specific lines of code. The gas costs of each trace is then aggregated and further processed to take into account refunds and external calls. Finally, a histogram of gas costs for each line of code across all executed traces is output. For programme counters that were never executed, only the static gas costs are calculated, but this may not be accurate or sound.

Next, in order generate the test inputs, a fuzzer by Ambroladze et al [49] is used. This generates transactions to run all public functions within a smart contract, and uses a feedbackloop to adjust arguments so as to achieve high code coverage. It also takes into account the timestamp of the deployment transaction, and fuzzes the block timestamps to execute functions that require a certain time to have passed.

Finally, to visualise the analysis, a webserver built on Flask was used, which links all three analysis components together.

3.2.2 Results

VisualGas reports a 88% average (and 94% median) code coverage for their fuzzer with a limit of 5,000 transactions, when testing against a dataset of 30,400 contracts

with less than 6,000 EVM instructions. They also report that larger contracts seemed to have lower coverage at that limit.

As for gas analysis, VisualGas did not evaluate their tool against a large dataset of contracts, or provide any measurement of precision for their analysis. Instead, they performed a case study on a TimedCrowdSale smart contract [50], and reported a running time of 45s and 82s for executing 5000 transactions and 10000 transactions respectively. They also claim a 99% code coverage for the SimpleToken contract, and a 95% coverage for the MyTimedCrowdSaleContract. However, no measurement of precision was provided for the contracts tested as well.

3.2.3 Limitations

Although VisualGas is closest to this project in terms of objectives, and aims to produce a gas cost estimate for every line of code, their implementation does not attempt to be sound or precise, considering the lack of evaluation for the precision of gas costs calculated. Rather, VisualGas at best serves to only provide a rough visualisation of which parts of code appear to be the most gas-consuming. In addition, their use of dynamic analysis in the form of programme traces results in a relatively average low code coverage of 88%, with a relatively high running time required to collect all traces. This is not ideal, since a large part of gas costs might still be hidden away within the code paths that were not traversed. Our implementation addresses this via symbolic execution, which reasons about a programme path-by-path rather than input-by-input, and can be significantly more efficient.

VisualGas also did not state their latest supported Solidity version, although images from their paper suggests this is 0.4.25, which was released more than 3 years ago. There is also no source code provided, and therefore we are unable to test their tool.

3.3 Summary of Research

From our findings, we discover that no current existing tool performs the objectives of this project well. GASTAP has a sophisticated static analysis framework and can notably derive parametric gas bounds in terms of function arguments, but only provides these at a transaction level rather than a per line-of-code level. VisualGas performs dynamic analysis to derive gas costs for each line of code, but has lower coverage and the precision of their estimates remains unknown.

Our project therefore aims to improve on these tools, by providing a gas cost estimate for each line of code, and enabling high code coverage via symbolic execution. We also plan to enhance the gas estimation of loops by performing static analysis on the Yul intermediate representation, to derive a parametric bound on the number of loop invocations.

4

Implementation

4.1 Feature summary

This project has successfully built and deployed Solbolt, a compiler explorer and gas analysis tool, which is available for all to use on the live website [51]. Next, we shall examine some of its key features, with reference to a summary graphic shown in Figure 4.2.

4.1.1 Solidity Compilation

The tool supports the compilation of contracts written in Solidity version 0.4.10 to 0.8.13, as well as support for EVM version from `homestead` to `berlin`. This is, to our knowledge, the most up-to-date and widest available compiler support for any Solidity gas analysis tool so far. The user also has the option to enable optimisation, as well as customise advanced options such as enabling the peephole optimiser, the inliner, the common subexpression eliminator, and more, as seen in Figure 4.1a. This may be helpful for developers or researchers investigating the effects of different optimisation options.

The tool also supports the compilation of contracts spread over multiple source files, as well as compiling multiple contracts from the same source. For verified contracts that are already deployed on chain, the tool also has Etherscan integration, and can directly load a contract from a verified Ethereum address for compilation, using Item 1.5 as indicated in the summary graphic.

After successful compilation of the source, the EVM assembly opcodes will be displayed on the right "Assembly viewer" pane. Sections of the opcodes will be colour-coded based on its mapping to the Solidity source code, similar to the Godbolt tool that inspired this project. Hovering over a section will also highlight and bring to view its corresponding mapped section on the other pane, which can be frozen via a keyboard shortcut.

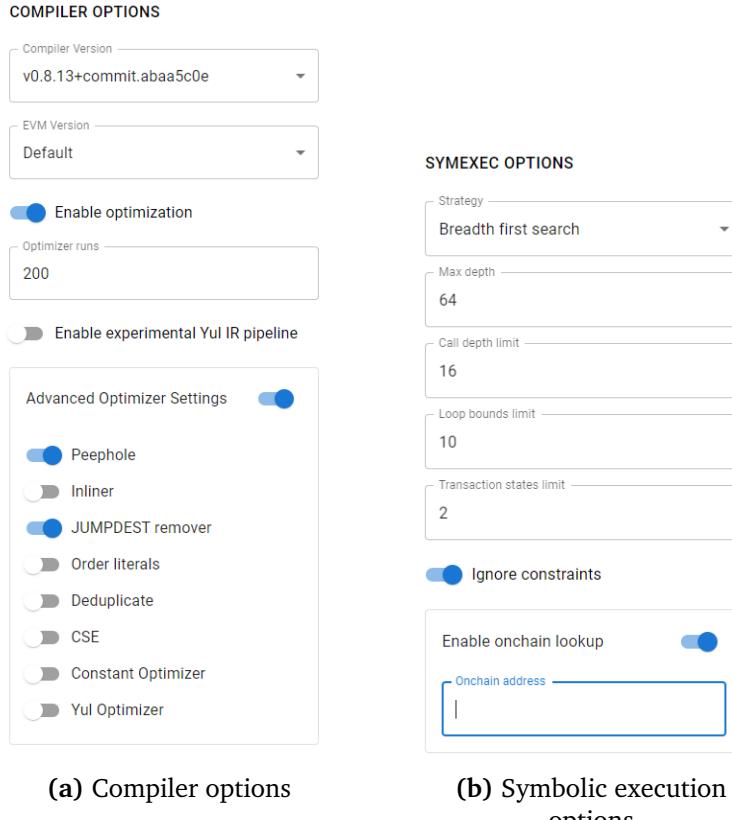


Figure 4.1: Solbolt options

4.1.2 Symbolic Execution

Once the source is compiled, users are able to run a symbolic execution using Item 2.3. The symbolic execution instance will try to execute each path of the contract bytecode, and derive gas estimates for each mapping generated by the compiler. This is indicated to the user as a coloured heatmap for quick identification of the gas-costly sections, with lighter yellow shades consuming fewer gas units, while the darker red shades consuming greater gas units.

Next, it will also calculate worst case gas estimates for each function it detects, as well as per-iteration gas costs for each loop, including hidden loops within compiler generated code. Finally, the symbolic executor can also detect common gas inefficient patterns, such as modifying a storage variable within a loop, which could possibly be fixed or refactored for additional gas savings.

Each symbolic execution is run with a timeout of 120 seconds, and users can customise some of its settings, such as the path traversal strategy, the maximum depth, the call depth limit, the loop bounds limit, and the number of transactions to run, as seen in Figure 4.1b.

This may be important to tweak for different contracts to achieve the best results, depending on how complex they are, how many times a loop is expected to run,

and so on. The user can also enable the Z3 SMT solver, which prunes paths that are not reachable as the constraints placed on them are not satisfiable. Users can also provide a on chain address for symbolic execution, and the engine will make use of the live concrete storage state of that address instead of a symbolic storage state.

4.1.3 Gas Inspector

Finally, the user is able to refer to the gas inspector tab on the right for a more detailed gas analysis report, as indicated by Item 3. Here, there is a gas heatmap legend indicating the total estimated gas used per transaction by the code section that is currently highlighted, as well as the gas used by the different categories, as seen with Item 3.4. It also shows the coverage percentage achieved by the symbolic execution, shown with Item 3.5.

The gas histogram, as seen with Item 3.6, shows a gas profile of the smart contract by plotting the total number of code sections that fall under a certain gas usage class. Lastly, Item 3.7 shows a histogram of worst case gas usage by each executed function.

4.2 Architecture Design

4.2.1 UI Elements and Redux State Management

The Solbolt frontend is built from scratch using React (TypeScript) 17.0.2 and Material UI 5.3.1. It is designed as a Single-Page Application (SPA), with three main panels for simple simple navigation, as seen in the screenshot in Figure 4.2. We chose to develop Solbolt fully in TypeScript rather than Javascript, because although this required the definition of additional interfaces and typing of arguments, it also provided benefits such as code hints and error highlighting when the expected types of objects did not match.

The frontend makes extensive use of redux, useContext and hooks. This prevented the common nuisance of passing props deeply into nested components when using React. For example, for managing the contract state, we created an ContractsContext that included the ContractsState, as well as functions to update the ContractsState, as seen in Listing 4.1.

```

1 const ContractsContext =
2   createContext<
3     [ContractsState | undefined, {
4       updateContract:
5         ((filename: string, name: string, contract: ContractJSON) => void) |
6           undefined,
7       updateAllContracts:
8         ((contracts: {[fileName: string]: {[contractName: string]: ContractJSON}}, |
9           ast: {[name: string]: any}) => void) | undefined
10    }]
11  >([undefined, {
12    updateContract: undefined,
13    updateAllContracts: undefined
14  }]
15 )
16 
```

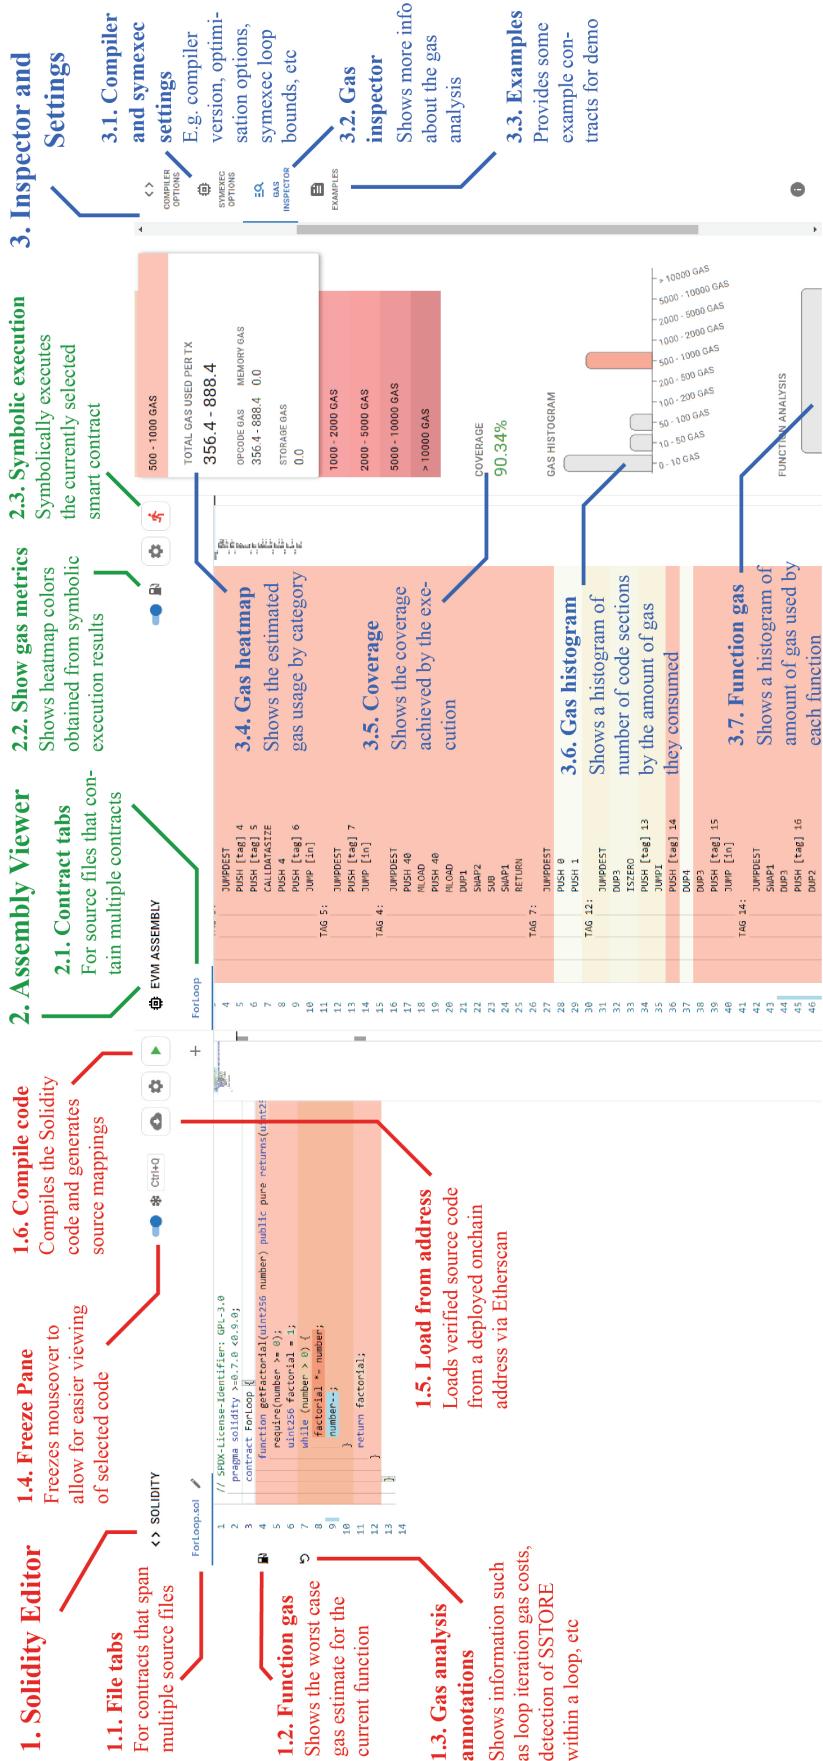


Figure 4.2: Summary of Solbolt features

```
13     }]);

```

Listing 4.1: ContractsContext used for managing contract state

ContractsContext also has a reducer, which takes an update type and a payload. Depending on the type of the update, the reducer will modify the state accordingly, making use of relevant data in the payload, as seen in Listing 4.2

```
1  function reducer(state: ContractsState, { type, payload }: { type: UpdateTypes,
2  	payload: PayloadType | undefined }) {
3  	switch (type) {
4  		case UpdateTypes.UPDATE_CONTRACT: {
5  			/* Perform update for one contract */
6  			...
7  		}
8
9  		case UpdateTypes.UPDATE_ALL_CONTRACTS: {
10  			/* Perform update for ALL contracts */
11  			...
12  		}
13
14  		default: {
15  			throw Error(`Unexpected action type in ContractsContext reducer: '${type}'`)
16  		}
17  }

```

Listing 4.2: Reducer used in ContractsContext

This reducer is called by the Provider, which lives near the root of the React DOM, above all components that wish to use the context. The class then exports hooks that can be easily used by any component that require the accessing or updating the contract state, such as `useContract`, `useUpdateContract`, and so on.

This method eliminates the need to pass props for shared information used frequently by many components across different routes, such as the application state, mappings state, highlighted class and contracts state. Updates are also handled automatically by React, because whenever the redux state is changed, the components dependent on it are also rerendered. This also makes it simple to store these states persistently in local browser storage, such that the user's session is restored after launching Solbolt again.

4.2.2 Monaco Editor

For the Solidity Editor and the EVM assembly viewer, we made use of the Monaco Editor, also used in Microsoft VSCode and the Godbolt compiler explorer. Although other similar libraries such as CodeMirror were lighter and had better mobile support, Monaco was the only editor that had native support for Solidity syntax highlighting, as well as a simpler system for handling code decorations and mouse handlers.

The current state of code decorations and mouse handlers are kept track of using a `useRef` variable, and destroyed when a new instance is being set to prevent memory leaks. Monaco also made use of a model and viewstate, which stores the content of

an editor as well as the current view and undo stack. These can be saved and loaded as needed when switching tabs.

4.2.3 Etherscan Integration

The Solbolt frontend is directly integrated with the Etherscan API, and has access to over 160,000 verified contract source codes to date [52]. Etherscan now also supports multiple source files for verification, and as such, Solbolt also introduced tabs support for editing and compiling multiple files. The user simply needs to input the Ethereum address they wish to analyse, and Solbolt will automatically send a request and parse the Etherscan result, before updating the view with the retrieved source files.

4.2.4 Backend and REST API

Requests to the Solbolt backend are first routed through the NGINX reverse-proxy, which also acts as a rate limiter. These are then forwarded to the API endpoints, built using Flask and served with Gunicorn. Each request for initiating a compilation or symbolic execution task will generate a new task ID, which is returned in the response. This task request is also simultaneously published to a Redis broker, which are being listened to by workers. The client keeps track of the task ID, and will poll for the task result periodically, until the task is completed and they are made available.

There is no model or database used in this project, and no Solidity code or symbolic execution results are stored permanently, other than being cached by Redis.

4.2.5 Celery Task Queue

For the actual execution of the published tasks, two Celery workers listen to tasks on the Redis broker, and pick them up as soon as they are ready and available. Each worker runs an instance of the extended Mythril symbolic execution engine or Solidity compiler, and once the job is complete, the results are stored within the Redis cache. These will be accessed and returned by the API endpoints upon the next client poll.

4.2.6 Deployment Pipeline

The frontend is built and served statically on AWS Amplify, which also utilises its global content distribution network for faster access. This is connected with the Github repository, and a new build is launched whenever a new commit is pushed.

The backend is instead hosted on a Hetzner Cloud instance with 3 AMD EPYC virtual cores and 4GB of RAM, running on Ubuntu 22.04 LTS. All backend services — the NGINX proxy, the Gunicorn instance, the Redis server, the Celery workers, and the SSL/TLS Certbot — are spun up at the same time using docker-compose, and

run within a Docker container. Deploying a new version is not automated, but requires manual pulling of the new Github commit, and then rebuilding the Docker containers.

Best practices for security are also observed at all times, such as placing all API keys and production secrets within environment variables. The backend is also protected from malicious bots and denial-of-service attacks via custom jails with fail2ban.

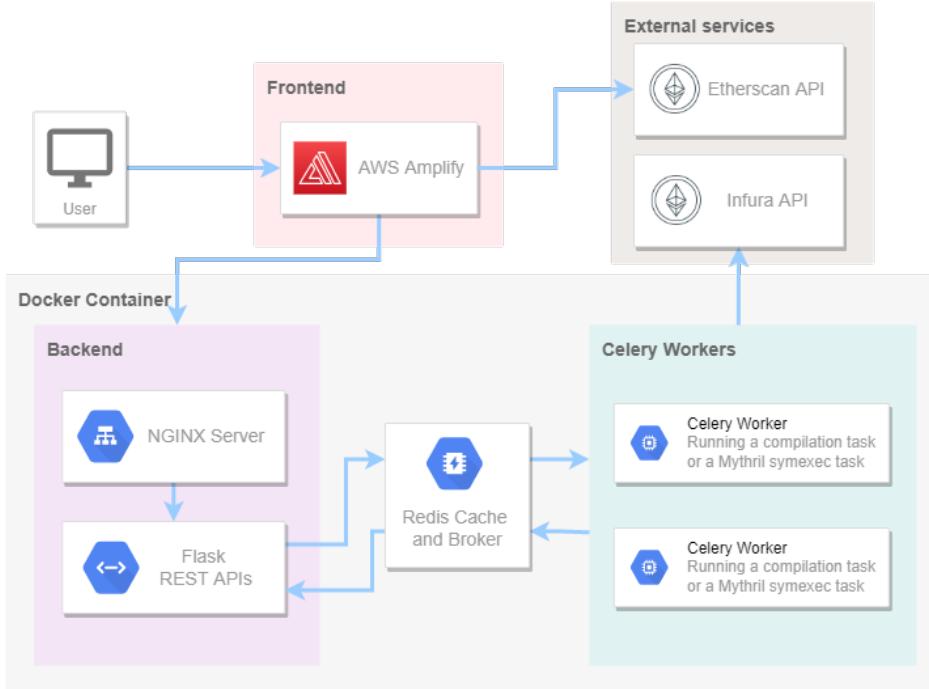


Figure 4.3: Solbolt architecture overview

4.3 Mytril Extensions

4.3.1 Gas Meter Plugin and Gas Hooks

The purpose of the gas meter plugin is to keep track of the average total amount of gas used per symbolic transaction by each Solidity code section. This will then be presented to the user in the form of a gas heatmap.

Originally, Mytril already has some gas tracking capabilities, but it only kept track of the total gas used up at each program point, and did not store any information about which sections of code used up how much of that gas, therefore having little utility for gas analysis. In addition, the original machine state counted SSTORE and SLOAD storage opcode gas costs together with the other opcodes, and therefore only stored the opcode and memory gas usage. However, because storage costs often represent the majority of gas usage in most smart contracts, we felt it was informative to separate out the storage costs from the rest of the opcode costs.

We therefore built a custom gas meter plugin, which stores a global gas meter object that in turn, stores a map of each code section to a struct containing the following information:

- `min_opcode_gas_used` and `max_opcode_gas_used` — stores the sum of the lower and upper bounds of the opcode gas estimated to be used
- `mem_gas_used` — stores the sum of the memory gas estimated to be used, as a result of memory expansion during execution
- `min_storage_gas_used` and `max_storage_gas_used` — stores the sum of the lower and upper bounds of the storage gas estimated to be used by `SSTORE` and `SLOAD` opcodes
- `num_invocations` — stores the total number of times that opcodes, which are mapped to this code section, are invoked throughout all transactions
- `num_tx` — stores the total number of transactions where opcodes mapped to this code section are executed

The plugin also uses a local gas meter object, which is stored in each global state as part of the gas meter annotation. This local gas meter contains the same information as above, but only for the current transaction trace.

In addition to the gas meter plugin, we also designed a new hook for the Mythril engine, called the gas hook. The pre-instruction hook was not suitable for this plugin, because the opcode has not been executed yet, and the gas is therefore not incremented. The instruction hooks are not suitable either, because they are called with the initial global state, rather than the new global states with updated gas costs. The post-instruction hooks are also not suitable, because although they had the updated gas costs after executing the opcode, the programme counter (PC) has also been incremented to the next instruction, and therefore we cannot attribute gas costs to the correct previous instruction that has used it. As such, we needed a new gas hook, which is called right after the machine state within the current global state is updated with the latest gas costs, but before the PC is incremented.

For every instruction, at the gas hook, the plugin runs the following pseudo-code in Listing 4.3. The gas plugin first gets the current gas meter annotation for the current global state, and then tries to obtain the source mapping for the current programme counter. If no source mapping is found, then the current code is compiler generated, and the plugin will attribute the gas cost of the current code to the last seen source mapping. Then, the plugin gets the difference in gas between the current machine state, and the gas last seen by the annotation (i.e. at the previous opcode). The plugin then updates the annotation by attributing the current gas difference to the source mapping, and updates the last seen gas of the annotation to the current gas.

```

1 class GasMeterPlugin:
2     @add_gas_hook(opcode = any):
3         function gas_meter_hook(state: GlobalState):
4             annotation = get_gas_meter_annotation(state)
5
6             pc = state.instruction["address"]
7

```

```

8     source_mapping = contract.get_source_mapping(pc)
9     if (contract.has_source(source_mapping.solidity_file_id)):
10        annotation.curr_key = parse_mapping(source_mapping)
11
12    if annotation.curr_key is not None:
13        gas_difference = get_difference_in_gas(annotation, state.mstate)
14
15        update_gas_meter_for_key(annotation.gas_meter, annotation.curr_key,
16                                   gas_difference)
17
18        set_last_seen_gas(annotation, gas_difference)
19
20    @add_pre_hook(opcode = STOP, REVERT, RETURN):
21    def end_hook(state: GlobalState):
22        annotation = get_gas_meter_annotation(state)
23
24        for key in annotation.gas_meter.keys():
25            merge(self.current_gas_meter[key], annotation.gas_meter[key])

```

Listing 4.3: Hooks implemented for gas meter plugin

When a transaction ending opcode is encountered, such as REVERT, STOP or RETURN, the local gas meter is merged into the global gas meter for each source mapping executed in the current ending transaction, which persists and stores this information across multiple transaction instances. After the symbolic execution is complete, statistics for each code section such as the mean are then calculated, and returned as part of the result to the client.

4.3.2 Function Tracker Plugin

The purpose of the function tracker plugin is to track the worst case transaction gas estimate for each function that can be called externally via a transaction. This is then displayed to the user via a glyph decoration within the Solidity editor, next to each function detected by the plugin.

Mythril also has some capabilities for tracking the current function that the execution is in. However, it only tracks the current active function, and therefore if function X calls another function Y of the same contract, and the transaction reverts in function Y , then the active function will still be function Y , and the worst case gas accumulated so far will be incorrectly attributed to function Y instead of function X . Instead, we want to know the entry function that is called at the start of the transaction, and attribute all gas accumulated throughout all possible paths (no matter if other functions are called as well) to that entry function.

The function tracker plugin was therefore created to store this information, and makes use of the EVM dispatcher routine to detect which function we have entered.

To understand how the plugin works, we first need to understand how the EVM dispatcher identifies and jumps to the correct function when executing a transaction call. Each smart contract call transaction must contain calldata, of which the starting 4 bytes indicate the signature (or hash) of the function that the transaction wishes to execute, as seen in Listing 4.4. At the entrypoint of the contract bytecode, the dispatcher first loads the calldata in line 2 of Listing 4.5, and then pushes a possible signature for a function in the smart contract, in line 5. Then, it compares the newly

pushed signature with that of the calldata using EQ, as seen in line 7, afterwhich it pushes the destination that the function resides in line 8, and jumps into it if the signatures are equal using JUMPI in line 9. This is then repeated for each function within the smart contract. If there are no matching functions found by the dispatcher, the transaction is reverted.

```

1 Transaction raw calldata:
2 0xa9059cbb // Function signature
3 00000....3f7614377 // Argument one
4 00000....11d800000 // Argument two

```

Listing 4.4: Transaction calldata

```

1 ...
2 CALLDATALOAD
3 DIV
4 AND
5 PUSH4 6FDDE03
6 DUP2
7 EQ
8 PUSH [tag] 2
9 JUMPI
10 DUP1
11 PUSH4 95EA7B3
12 EQ
13 PUSH [tag] 3
14 JUMPI
15 DUP1
16 PUSH4 18160DDD
17 ...

```

Listing 4.5: EVM dispatcher routine

Therefore, we can use this pattern to identify which function the current trace entered at the beginning. Since the calldata is symbolic, we cannot simply check its value for the function signature called, as this value does not yet exist. Instead, we make use of the algorithm outlined by the pseudocode in Listing 4.6.

```

1 class FunctionTrackerPlugin:
2     @add_pre_hook(opcode = PUSH4):
3         function push_hook(state: GlobalState):
4             if not isinstance(state.current_transaction, ContractCreationTransaction):
5                 annotation = get_function_tracker_annotation(state)
6
7                 push_value = state.instruction["argument"]
8
9                 if (annotation.current_function == None):
10                     fn_signature = parse_signature(push_value)
11
12                     if (fn_signature in contract.signatures and state.code[state.mstate.pc +
13                         1]["opcode"] == "EQ"):
14                         fn_name = contract.signatures[fn_signature]
15                         annotation.candidate_function = f'{parsed_value}:{fn_name}'
16
17             @add_post_hook(opcode = JUMPI):
18                 function jump_hook(state: GlobalState):
19                     if not isinstance(state.current_transaction, ContractCreationTransaction):
20                         annotation = get_function_tracker_annotation(state)
21
22                     if annotation.current_function == None and annotation.candidate_function != None:
23
24                         prev_pc = state.mstate.prev_pc
25                         curr_pc = state.mstate.pc
26
27                         did_jump_into_function = prev_pc + 1 != curr_pc
28
29                         if (did_jump_into_function):
30                             annotation.current_function = annotation.candidate_function
31
32                         annotation.candidate_function = None

```

```

33     @add_pre_hook(opcode = STOP, REVERT, RETURN):
34         function end_hook(state: GlobalState):
35             annotation = get_function_tracker_annotation(state)
36
37             if (annotation.current_function != None):
38                 prev_max_gas = global_function_gas_meter.get(annotation.current_function, 0)
39                 global_function_gas_meter[annotation.current_function] = max(prev_max_gas,
40                     state.mstate.max_gas_used)

```

Listing 4.6: Hooks implemented for function tracker plugin

First, for both the JUMPI and PUSH4 hooks, we check if the current transaction is the contract creation transaction. Since we only want to collect data for the runtime transactions, we do nothing if this is true. Otherwise, we first look for a PUSH4 instruction with an argument that matches one of the function signatures in the contract. We also check if the next instruction is an EQ opcode. If this matches, then it is likely that we are in the dispatcher routine, and we set the current candidate function to be this function signature.

After that, we look for the upcoming JUMPI instruction, as per the routine. Since we are using the post hook for this, the global state provided would be the resulting state after the JUMPI has already been executed. Therefore, we can check if the state has jumped into the function (the true case), or simply moved on to the next instruction in the dispatcher routine (the false case), by checking if the current PC is equal to the previous PC plus one. In the true case, we simply mark the current function in the annotation as the candidate function (which we just jumped into). In the false case, we set the current function back to None, and repeat the process again for the next function signatures.

Once the current function has been identified, the plugin will not for this dispatcher routine anymore. This is to ensure that the original entry function (the one that was called in a transaction) will not be overriden by another function that might be called within the entry function.

When a transaction ending opcode is encountered, we then compare the total gas accumulated in the current global state, with the maximum gas encountered so far by the global function gas meter, for the current function. We finally store the new maximum of both gas counts for that function, and repeat this over multiple transactions. The final maximum gas counts for each function is then returned as part of the result to the client.

4.3.3 Loop Gas Meter Plugin

Next, we developed a loop gas meter plugin, aimed to keep track of per-iteration loop gas costs for each loop that was detected. The Mytril engine already has a bounded loops strategy, which is able to detect the current number of loop iterations, but we would like to keep track of the gas cost associated with this as well.

To do this, we make use of the following pseudocode in Listing 4.7, as follows.

```

1  class FunctionTrackerPlugin:
2      @add_pre_hook(opcode = any):

```

```

3     function pre_instr_hook(state: GlobalState):
4         if not isinstance(state.current_transaction, ContractCreationTransaction):
5             annotation = get_loop_gas_meter_annotation(state)
6
7             pc = state.instruction["address"]
8             source_mapping = get_source_mapping(pc)
9
10            if source_mapping is not None:
11                annotation.curr_key = source_mapping
12
13    @add_pre_hook(opcode = JUMPDEST):
14        function jumpdest_hook(state: GlobalState):
15            annotation = get_loop_gas_meter_annotation(state)
16
17            pc = state.instruction["address"]
18
19            new_trace_item = TraceItem(pc, state.mstate.max_gas_used)
20            annotation.trace.append(new_trace_item)
21
22            (loop_head, loop_gas) = find_loop(annotation.trace)
23
24            if loop_head is not None:
25                source_mapping = contract.get_source_mapping(loop_head, constructor=self.
26                    is_creation)
27
28                is_hidden_loop = is_generated_code(source_mapping)
29
30                append_gas_to_gas_meter(annotation, loop_head, loop_gas, is_hidden_loop)
31
32    @add_pre_hook(opcode = STOP, REVERT, RETURN):
33        function end_hook(state: GlobalState):
34            annotation = get_loop_gas_meter_annotation(state)
            merge(global_loop_gas_meter, annotation.loop_gas_meter)

```

Listing 4.7: Hooks implemented for loop gas meter plugin

First, we set up a pre-instruction hook for all opcodes. This checks if the current transaction is not a contract creation transaction, and then gets the current PC and source mapping of the PC. If the source mapping is not none, then we update the current source mapping for the annotation. This essentially updates the annotation with the latest source mapping reached at every point of the execution.

Next, we have a pre-hook for all JUMPDEST opcodes. These opcodes mark the start of every valid jump destination, which the loop body would start with. We first append the annotation trace (which keeps a log of the PCs of all JUMPDEST opcodes seen so far) with the current PC, and then call the `find_loop` function to try to find the PC of the current loop head, as well as the current loop iteration gas cost. We will explain in more detail about how this function works in the next section. If no loop is found, the loop head returned is `None`. If a loop is found, we then check if the current loop is a hidden one (where it is generated by the compiler, and the source mapping does not exist). Finally, all of this information is added to the local loop gas meter within the current annotation.

Lastly, we have a pre-hook for all terminating opcodes. Once these opcodes are reached, we get the current annotation (and hence local loop gas meter) from the global state, and merge it with the global loop gas meter stored in the plugin. The average iteration gas cost for each loop head found is then taken, and returned to the user as the result.

Now, we explain more about how the `find_loop` function works. An outline of its pseudocode is provided in Listing 4.8. Essentially, the function works by finding a match of the latest two consecutive JUMPDEST instructions from within the PC history stored in the trace. If this is found, we set as former of the latest two consecutive instructions as the loop head. This heuristic is not 100% accurate however, and we may find false positives. For example, if two sections of a function both call the same helper function, this may register as a loop even though the helper function is not recursive. However, we chose to use this heuristic over other static analysis methods (on the EVM bytecode or Solidity code) because it is able to not only capture possible "hidden" loops generated by the compiler, it is also able to capture any recursive loops seen during each symbolic execution. Also, from our testing, the rate of false positives is not high, and can be trivially ignored by the developer.

```

1  function find_loop(trace):
2      found_loop_head = None
3
4      loop_head_index = 0
5
6      for loop_head_index in range(len(trace) - 3, -1, -1):
7          if trace[loop_head_index].pc == trace[-2].pc and trace[loop_head_index + 1].pc
8              == trace[-1].pc:
9                  found_loop_head = trace[loop_head_index]
10                 break
11
12     if found_loop_head:
13         found_pc = found_loop_head.pc
14         found_gas = found_loop_head.gas
15
16         loop_iteration_gas = trace[-2].gas - found_gas
17
18         // remove loop from trace to prevent finding it again next time
19         for i in range(len(trace) - 3, loop_head_index - 1, -1):
20             del trace[i]
21
22     return (found_pc, loop_iteration_gas)
23 else:
24     return (None, 0)
```

Listing 4.8: `find_loop` function

To illustrate this, let us examine Figure 4.4. Here, we begin at PC 241, which is the first PC to be added to the previously empty trace. We next jump to PC 388, which for this example is the loop head containing the break condition. This is also appended to the trace. Next, we jump to PC 411, the loop body, and append it to the trace as well. This takes us back to PC 388, which is appended to the trace again. Currently, the trace contains {241, 388, 411, 388}, which suggests there might be a loop, but because no two consecutive PCs are found yet, the algorithm does not determine a loop exists so far. We now move to PC 411 again in the second iteration of the loop, and now the trace contains {241, 388, 411, 388, 411}. Here, the algorithm finally finds a match, and correctly determines a loop starting at PC 388. It then calculates the gas difference between the first and second instances of arriving at PC 388, which is the gas cost of one iteration.

After this, we delete all trace items seen in the interval between the previous loop head and the current PC. In our example, in step 6, we can see that the second occurrence of {388, 411} has been removed. This is to prevent any double counting of

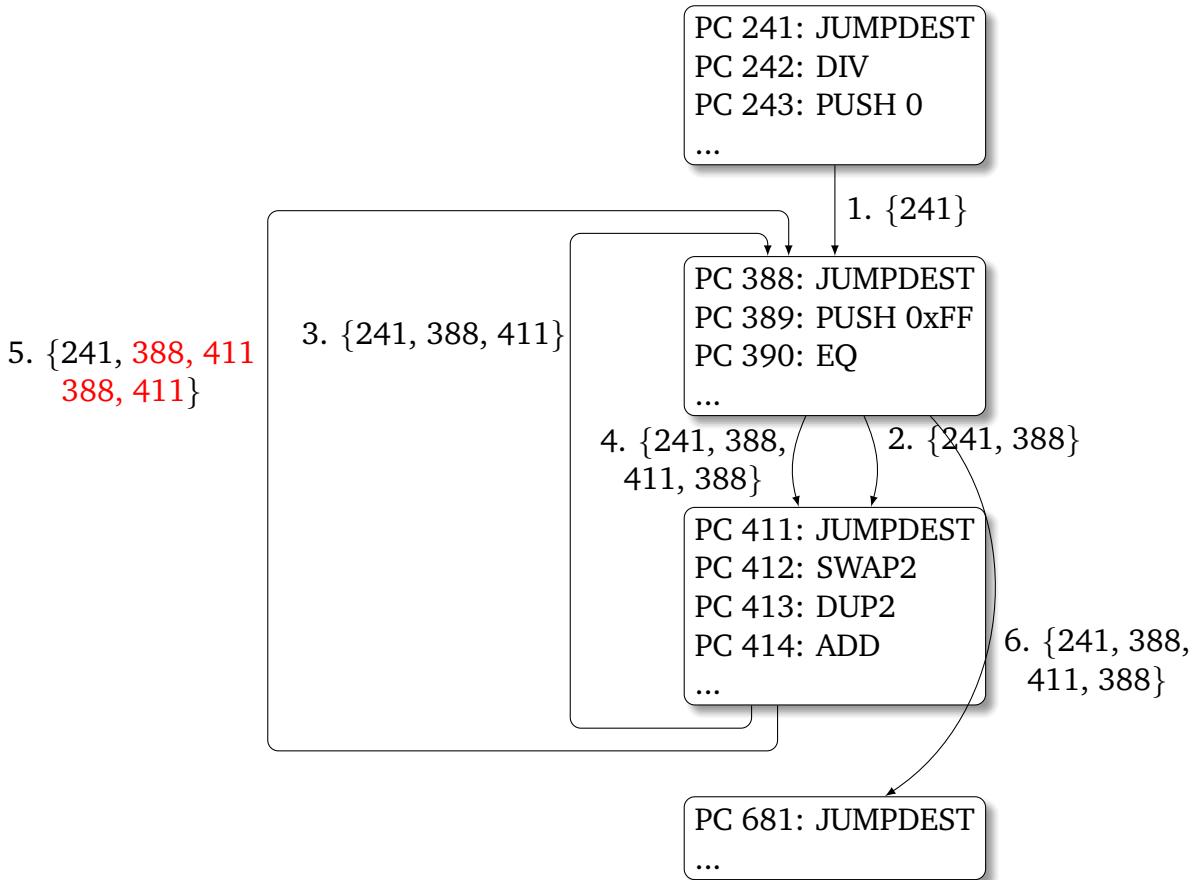


Figure 4.4: Example of loop finding algorithm

loop heads. For example, if no trace items were deleted, the trace at step 6 would instead be {241, 388, 411, 388, 411, 388}, and the algorithm would detect another loop at PC 411. However, this "loop" was already accounted for in the loop starting at PC 388. Finally, the PC of the loop head, and the iteration gas cost is returned.

4.3.4 Loop Mutation Detector Plugin

Finally, we also developed a loop mutation detector, which checks if there are any SLOAD or SSTORE opcodes found within a loop. This is because one of the common gas-hungry anti-patterns is unnecessarily modifying storage variables within the loop body, when you can instead use a temporary local copy of the storage variable within the loop, and then update the actual storage variable just once when the loop is complete. This example is evaluated in more detail in Section TODO: Add gas hungry loop example.

To implement this plugin, we set up an instruction hook for SSTORE and SLOAD instructions, and simply piggy-backed off the `JumpdestCountAnnotation` used by the bounded loops strategy to determine if we are in a loop or not. If we are, then the plugin looks up the source mapping for the current PC, and then adds this to the dictionary storing the set of offending source mappings.

The plugin is not always accurate, and discretion is required from the developer to determine whether the code can be refactored to fix this. For example, in the case of updating dynamically-sized variable types, such as `string` and `bytes`, the compiler automatically generates a routine that involves iterating through and updating the relevant storage slots in a loop. This will be flagged by the plugin, but cannot be fixed or refactored, firstly because the offending code is compiler generated, and secondly because this actually requires the use of a loop to iterate through the different slots.

4.3.5 Removal of Z3 SMT Constraints

While testing the performance of the symbolic execution engine, we noticed that the first few transactions were completed very quickly, but the next transactions took exponentially longer times to symbolically execute. This may be because the state space grows exponentially with the number of transactions, but we also suspected that because the list of constraints for each state also grows exponentially, this makes solving them much slower. In addition, we found that some states were not being executed, because the Z3 solver was timing out and simply marking the state as unreachable.

Security verification requires the fact that unreachable states are pruned and not executed, because they are focused on which paths are reachable and which are not. However, for our purpose of gas estimation, this is not as important as faster execution times and greater coverage. Being able to quickly calculate the gas usage of a code section when it is reached is more important to us than knowing exactly if that particular section is reachable or not. As such, although useful for pruning away states when the list of constraints is relatively small, the cost of running this Z3 solver quickly outweighs its benefits when the list of constraints starts to grow exponentially.

Therefore, we provide the user with a setting that disables the Z3 constraint check, and simply runs every possible execution path regardless of whether they are reachable or not. This is turned on by default, and in our testing, provides much better execution coverage with for the same execution time.

5

Evaluation

In order to evaluate the quality and effectiveness of the tool created from the project, we plan to investigate the following categories:

1. **Speed**, in terms of the median and average time needed to analyse a contract
2. **Coverage**, in terms of the percentage of EVM bytecode for which the gas estimates can be derived
3. **Precision**, in terms of the gas estimates calculated
4. **User experience**, in terms of the features of the tool and how easy they are to use

For items 1 and 2, we aim to evaluate our tool against a dataset of over 70,000 unique Ethereum smart contracts [52], and collect statistics regarding run time and code coverage. These will be compared with other existing tools such as GASTAP and VisualGas. The tools are not open source, and therefore we are only able to compare with the published statistics, and are unable to verify the results ourselves in a similar setting.

Then, for item 3, we aim to evaluate the gas bounds derived from our tool against the gas cost of real transactions, using 100 real transactions from 100 contracts randomly chosen from the dataset. Then, we can compare the precision of these gas bounds with those generated by the Solidity compiler. We can also compare the precision in terms of overhead with GASTAP. These can only be evaluated at a transaction level, since no other existing tool is able to calculate precise gas bounds at a line of code level.

Finally, for item 4, we aim to evaluate the features of our tool against other existing tools, such as the Remix IDE and VisualGas. We also plan on gathering user feedback based on the following metrics:

1. **Responsiveness**, or the total time wasted on waiting as a percentage of total time spent on performing a task

2. **System Usability Scale**, which is a questionnaire of ten items for rating overall user experience
3. **Error Occurrence Rate**, which is the number of times an error occurs for each task performed
4. **Average Time on Task**, for the first attempt as well as for repeat attempts
5. **Hint rate**, which is the number of times a user has to ask for hints on how to perform a particular task

We plan on testing the tool with real users to gather these metrics, by asking them to perform the following tasks:

- Import a Solidity project
- Generate the EVM code mappings for the Solidity code
- Generate the Yul code mappings for the Solidity code
- Generate a heatmap of gas costs for the Solidity code

To consider the project a success, we would ideally like to achieve a speed comparable or better than existing gas analysis tools, a higher coverage than existing tools, and comparable or better precision than existing tools. This is because most existing tools only evaluate gas costs at a transaction level, and breaking it down into line of code level may incur some loss of precision. We would also like to create a rich user experience that has a high usability score, and is therefore simple to navigate but powerful to use.

6

Ethical considerations

Overall, there does not seem to be any major ethical concerns regarding the nature of this project. We have also investigated the licenses for the libraries we intend to incorporate within our tool, namely Mythril and the Solidity compiler, which respectively have the MIT and GPL-3.0 licenses, and should be permissive enough for our nature of work. However, there are still certain key points of consideration. First of all, because the project deals with decompiling Solidity code into its Yul intermediate representation and EVM code, and generating the mappings between them, they may reveal or expose vulnerabilities within the compiler generated code. With Ethereum smart contracts nowadays holding up to millions of dollars worth of funds, if such information is revealed to malicious actors, it could result in a permanent and irrevocable loss of funds.

In addition, this feature may also allow malicious actors to more easily obfuscate their code, by possibly implanting fragments of malicious inline assembly into template contract codes, claiming they serve another purpose. Then, users who do not analyse the code in detail and simply reuse them may at best lead to bugs in execution, or at worst loss of funds.

Next, there is no guarantee provided by our project about the precision or reliability of the gas analysis generated. Therefore, if the gas analysis is not accurate, and other projects do not carefully consider the gas used by their contract calls and simply rely on estimates derived from our tools, it may result in out of gas errors in future transactions, which may in turn lead to funds being locked forever and lost. We therefore intend to waive any liabilities derived from the use of our tool.

Also, generating the Yul intermediate representation is still an experimental function of the Solidity compiler that we make use of. Therefore, it is not assured to be completely correct, and users should not use the bytecode generated by our code directly when deploying contracts.

Finally, when collecting user data during the evaluation of our project, we must also consider the potential privacy risks and the right to withdraw from the survey. Users will not be asked to provide any personal information, and the data collected will be

anonymised and aggregated, to prevent any leaks of sensitive information.

7

Project Plan

As of the writing of this report, the following foundational steps have been completed:

- Testing of the Solidity compiler for its capabilities and scope of debug logs emitted
- Testing of various symbolic execution engines (Mythril and Oyente) to investigate their capabilities
- Creation of a basic web application that is able to decompile a Solidity contract using the javascript Solidity compiler
- Research on previous gas estimation techniques for EVM, as well as various loop bounds analysis methods for symbolic execution

Moving forward, the plan for the project is as follows:

1. Building the compilation pipeline for the basic web app that automatically compiles Solidity code into both EVM and Yul IR. **Suggested timeline: 1.5 weeks (31 Jan - 9 Feb)**
2. Development of the mapping feature that allows compiled EVM and Yul code to be mapped to the corresponding Solidity code. **Suggested timeline: 1.5 weeks (10 Feb - 21 Feb)**
3. Extending the Mythril symbolic execution engine to provide aggregate gas estimates for each basic block. **Suggested timeline: 2 weeks (22 Feb - 7 Mar)**
4. Development of the heatmap feature that takes the gas estimates calculated and displays it in a heat map, on top of the compiled EVM and Yul IR. **Suggested timeline: 2 weeks (7 Mar - 21 Mar)**
5. Extension of the Solidity Yul compiler with our loop analysis algorithm, to emit parametric loop bounds within the output EVM. **Suggested timeline: 3 weeks (21 Mar - 11 Apr)**

6. Extension of web app to take into account the emitted parametric loop bounds during gas estimate calculation. **Suggested timeline: 1.5 weeks (14 Apr - 21 Apr)**
7. (*Stretch*) Development of a symbolic debugger using the Mythril engine, that allows the user to step through an execution and examine the values within the stack, the programme counter, and the gas used. **Suggested timeline: 4 weeks**
8. (*Fallback*) Development of a feature that builds and displays the control flow graph for both Yul and EVM bytecode **Suggested timeline: 2 weeks**
9. Conduct evaluation on the web application, and collect user feedback **Suggested timeline: 2 weeks (10 May - 24 May)**
10. Writing of draft report, and meet with supervisor **Suggested timeline: 2 weeks (24 May - 7 Jun)**
11. Finalise report and create and rehearse for final presentation **Suggested timeline: 2 weeks (7 Jun - 20 Jun)**

For items 5 and 6, if it is found to be unable to achieve the stated goal in time, or if the previous steps take longer than expected, the project will fallback into developing step 8 instead, which should be more manageable to complete. If the project goes very well and there is extra time left, we will focus on developing step 7, which would be an additional (but useful) feature for the tool. We have also included 2 weeks of buffer time for courseworks and examinations.

Bibliography

- [1] Yoav Vilner. Some rights and wrongs about blockchain for the holiday season, 2018. URL <https://www.forbes.com/sites/yoavvilner/2018/12/13/some-rights-and-wrongs-about-blockchain-for-the-holiday-season/>. pages 1
- [2] Hayden Adams, Noah Zinsmeister, Moody Salem, River Keefer, Dan Robinson. Uniswap v3 core. URL <https://uniswap.org/whitepaper-v3.pdf>. pages 1
- [3] Aave - protocol whitepaper. URL <https://github.com/aave/protocol-v2/blob/master/aave-v2-whitepaper.pdf>. pages 1
- [4] Robert Leshner, Geoffrey Hayes. Compound: The money market protocol. URL <https://compound.finance/documents/Compound.Whitepaper.pdf>. pages 1
- [5] MakerDAO. The maker protocol: Makerdao's multi-collateral dai (mcd) system. URL <https://makerdao.com/en/whitepaper/>. pages 1
- [6] Evan Kereiakes, Do Kwon, Marco Di Maggio, Nicholas Platias. Terra money: Stability and adoption. URL https://assets.website-files.com/611153e7af981472d8da199c/618b02d13e938ae1f8ad1e45_Terra_Whitepaper.pdf. pages 1
- [7] Axie infinity whitepaper. URL <https://whitepaper.axieinfinity.com/>. pages 1
- [8] The sandbox whitepaper. URL https://installers.sandbox.game/The_Sandbox_Whitepaper_2020.pdf. pages 1
- [9] Defi llama: Ethereum tvl. URL <https://defillama.com/chain/Ethereum>. pages 2
- [10] Ethereum average gas price chart. URL <https://etherscan.io/chart/gasprice>. pages 2
- [11] Harry Robertson. Ethereum transaction fees are running sky-high. that's infuriating users and boosting rivals like solana and avalanche. URL <https://markets.businessinsider.com/news/currencies/ethereum-transaction-gas-fees-high-solana-avalanche-cardano-crypto-blockchain-2>. pages 2

- [12] Ishan Pandey. Ethereum’s high gas fees is limiting on-chain activities. URL <https://hackernoon.com/ethereums-high-gas-fees-is-limiting-on-chain-activities>. pages 2
- [13] Samuel Wan. Uniswap activity sends ethereum gas fees sky high. URL <https://www.newsbtc.com/news/ethereum/uniswap-activity-sends-ethereum-gas-fees-sky-high/>. pages 2
- [14] Monika di Angelo and Gernot Salzer. A survey of tools for analyzing ethereum smart contracts. In *2019 IEEE International Conference on Decentralized Applications and Infrastructures (DAPPCon)*, pages 69–78, 2019. doi: 10.1109/DAPPCon.2019.00018. pages 2
- [15] *Solidity Documentation Release 0.8.12*, 2021. URL <https://buildmedia.readthedocs.org/media/pdf/solidity/develop/solidity.pdf>. pages 2, 8
- [16] *Solidity v0.7.0 Breaking Changes*, 2021. URL <https://docs.soliditylang.org/en/v0.7.0/070-breaking-changes.html>. pages 2
- [17] *Solidity v0.8.0 Breaking Changes*, 2021. URL <https://docs.soliditylang.org/en/v0.8.10/080-breaking-changes.html>. pages 2
- [18] ConsenSys. *A Definitive List of Ethereum Developer Tools*, 2021. URL <https://media.consensys.net/an-definitive-list-of-ethereum-developer-tools-2159ce865974>. pages 2
- [19] Ting Chen, Youzheng Feng, Zihao Li, Hao Zhou, Xiaopu Luo, Xiaoqi Li, Xiuzhuo Xiao, Jiachi Chen, and Xiaosong Zhang. Gaschecker: Scalable analysis for discovering gas-inefficient smart contracts. *IEEE Transactions on Emerging Topics in Computing*, 9(3):1433–1448, 2021. doi: 10.1109/TETC.2020.2979019. pages 2
- [20] Ethererik. Governmental’s 1100 eth jackpot payout is stuck because it uses too much gas, 2016. URL https://www.reddit.com/r/ethereum/comments/4ghzhv/governmentals_1100_eth_jackpot_payout_is_stuck/. pages 2
- [21] Bitcoin Suisse. Compatible competition - empowering or encroaching on ethereum?, 2016. URL <https://www.bitcoinsuisse.com/research/decrypt/compatible-competition-empowering-or-encroaching-on-ethereum>. pages 3
- [22] Matt Godbolt. Optimizations in c++ compilers. *Commun. ACM*, 63(2):41–49, jan 2020. ISSN 0001-0782. doi: 10.1145/3369754. URL <https://doi.org/10.1145/3369754>. pages 3

- [23] Gavin Wood. Ethereum: A secure decentralised generalised transaction ledger. Dec 2021. URL <https://ethereum.github.io/yellowpaper/paper.pdf>. pages 3, 5, 6, 7, 8
- [24] Ainur R. Zamanov, Vladimir A. Erokhin, and Pavel S. Fedotov. Asic-resistant hash functions. In *2018 IEEE Conference of Russian Young Researchers in Electrical and Electronic Engineering (EIConRus)*, pages 394–396, 2018. doi: 10.1109/EIConRus.2018.8317115. pages 5
- [25] Ethereum. Proof-of-work (pow), . URL <https://ethereum.org/en/developers/docs/consensus-mechanisms/pow/>. pages 5
- [26] Nick Szabo. Smart contracts, 1994. URL <https://www.fon.hum.uva.nl/rob/Courses/InformationInSpeech/CDROM/Literature/LOTwinterschool2006/szabo.best.vwh.net/smart.contracts.html>. pages 5
- [27] Nick Szabo. Formalizing and securing relationships on public networks. *First Monday*, 2(9), Sep. 1997. doi: 10.5210/fm.v2i9.548. URL <https://firstmonday.org/ojs/index.php/fm/article/view/548>. pages 6
- [28] Konstantinos Christidis and Michael Devetsikiotis. Blockchains and smart contracts for the internet of things. *IEEE Access*, 4:2292–2303, 2016. doi: 10.1109/ACCESS.2016.2566339. pages 6
- [29] Ethereum. Ethereum virtual machine (evm), . URL <https://ethereum.org/en/developers/docs/evm/>. pages 6
- [30] Solidity. *Introduction to Smart Contracts*, 2021. URL <https://docs.soliditylang.org/en/latest/introduction-to-smart-contracts.html>. pages 6, 7
- [31] Juan Benet. IPFS - content addressed, versioned, P2P file system. *CoRR*, abs/1407.3561, 2014. URL <http://arxiv.org/abs/1407.3561>. pages 6
- [32] Samuel Williams, William Jones. Arweave lightpaper, Apr. 2018. URL <https://www.arweave.org/files/arweave-lightpaper.pdf>. pages 6
- [33] Ethereum. Gas and fees, . URL <https://ethereum.org/en/developers/docs/gas/>. pages 7
- [34] Vitalik Buterin, Eric Conner, Rick Dudley, Matthew Slipper, Ian Norden, Abdelhamid Bakhta. Eip-1559: Fee market change for eth 1.0 chain. Apr 2019. URL <https://eips.ethereum.org/EIPS/eip-1559>. pages 7
- [35] Ethereum. Design rationale, . URL <https://eth.wiki/en/fundamentals/design-rationale>. pages 7
- [36] Ethereum. Smart contract languages, . URL <https://ethereum.org/en/developers/docs/smart-contracts/languages/>. pages 8

- [37] Fabian Vogelsteller, Vitalik Buterin. Eip-20: Token standard. URL <https://eips.ethereum.org/EIPS/eip-20>. pages 8
- [38] William Entriken, Dieter Shirley, Jacob Evans, Nastassia Sachs. Eip-721: Non-fungible token standard. URL <https://eips.ethereum.org/EIPS/eip-721>. pages 8
- [39] OpenZeppelin. Openzeppelin contracts. URL <https://github.com/OpenZeppelin/openzeppelin-contracts>. pages 9
- [40] Solidity. Expressions and control structures, . URL <https://docs.soliditylang.org/en/latest/control-structures.html>. pages 9
- [41] Solidity. Yul, . URL <https://docs.soliditylang.org/en/latest/yul.html>. pages 10
- [42] James C. King. Symbolic execution and program testing. *Commun. ACM*, 19(7):385–394, jul 1976. ISSN 0001-0782. doi: 10.1145/360248.360252. URL <https://doi.org/10.1145/360248.360252>. pages 11
- [43] Christopher Signer. Gas cost analysis for ethereum smart contracts. 2018. URL https://www.research-collection.ethz.ch/bitstream/handle/20.500.11850/312914/TMeineOrdnerMaster-ArbeitenHS18Signer_Christopher.pdf. pages 11, 17
- [44] Loi Luu, Duc-Hiep Chu, Hrishi Olickel, Prateek Saxena, and Aquinas Hobor. Making smart contracts smarter. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, CCS ’16, page 254–269, New York, NY, USA, 2016. Association for Computing Machinery. ISBN 9781450341394. doi: 10.1145/2976749.2978309. URL <https://doi.org/10.1145/2976749.2978309>. pages 12
- [45] Mark Mossberg, Felipe Manzano, Eric Hennenfent, Alex Groce, Gustavo Grieco, Josselin Feist, Trent Brunson, and Artem Dinaburg. Manticore: A user-friendly symbolic execution framework for binaries and smart contracts. *CoRR*, abs/1907.03890, 2019. URL <http://arxiv.org/abs/1907.03890>. pages 12
- [46] Bernhard Mueller. Smashing ethereum contracts for fun and real profit. *HITBSecConf*, 2018. URL <https://github.com/b-mueller/smashing-smart-contracts/>. pages 12
- [47] Elvira Albert, Jesús Correas, Pablo Gordillo, Guillermo Román Díez, and Albert Rubio. Don’t run on fumes — parametric gas bounds for smart contracts. *Journal of Systems and Software*, 176:110923, 02 2021. doi: 10.1016/j.jss.2021.110923. pages 15
- [48] Elvira Albert, Jesús Correas, Pablo Gordillo, Guillermo Román-Díez, and Albert Rubio. Analyzing smart contracts: From EVM to a sound control-flow graph. *CoRR*, abs/2004.14437, 2020. URL <https://arxiv.org/abs/2004.14437>. pages 16

- [49] Mislav Balunović, Nodar Ambroladze, Petar Tsankov, and Martin Vechev. Learning to fuzz from symbolic execution with application to smart contracts. pages 531–548, 11 2019. ISBN 978-1-4503-6747-9. doi: 10.1145/3319535.3363230. pages 17
- [50] Christopher Signer. Timedcrowdsale. URL <https://github.com/ABBDVD/TimedCrowdsale>. pages 18
- [51] Richard Xiong. Solbolt. URL <https://www.solbolt.com>. pages 19
- [52] Martin Ortner and Shayan Eskandari. Smart contract sanctuary. URL <https://github.com/tintinweb/smarty-contract-sanctuary>. pages 24, 34